

```

import os
import tempfile
import matplotlib as mpl
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from termcolor import colored as cl
import itertools
import tensorflow as tf
from tensorflow import keras
import seaborn as sns
import sklearn
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

mpl.rcParams['figure.figsize'] = (12, 10)
colors = plt.rcParams['axes.prop_cycle'].by_key()['color']

```

```
!curl -L -O "https://training-images-4995.s3.amazonaws.com/data.zip"
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
			Dload Upload	Total	Spent	Left	Speed
100 69.0M	100 69.0M	0 0	28.7M 0	0:00:02	0:00:02	--:--:--	28.7M

```

from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.metrics import f1_score

```

```
!unzip -q "data.zip" -d .
```

```
!head -n 3 processed_creditcards.csv
```

```

,V1,V2,V3,V4,V5,V6,V7,V8,V9,V10,V11,V12,V13,V14,V15,V16,V17,V18,V19,V20,V21,V22,
0,-1.3598071336738,-0.0727811733098497,2.53634673796914,1.37815522427443,-0.3383
1,1.1918571113148602,0.26615071205963,0.16648011335321,0.448154078460911,0.06001

```

```

from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier

```

```
df = pd.read_csv('processed_creditcards.csv')
df.head()
```

	Unnamed: 0	V1	V2	V3	V4	V5	V6	V7	
0	0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.
1	1	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.
2	2	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.
3	3	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.
4	4	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.

```
total_cases = len(df)
valid_transaction_count = len(df[df.Class == 0])
invalid_transaction_count = len(df[df.Class == 1])
perc_fraudulent = round(fraud_count/nonfraud_count*100, 2)
```

```
print("There are", total_cases, "Total Cases")
print("Total", valid_transaction_count, "Not frauduelnt cases")
print("Total", invalid_transaction_count, "actual frauduelnt cases")
print("RESULT = 100*(", invalid_transaction_count, "/", total_cases, ") = ", 100*inva
```

```
There are 284807 Total Cases
Total 284315 Not frauduelnt cases
Total 492 actual frauduelnt cases
RESULT = 100*( 492 / 284807 ) = 0.1727485630620034 PERCENT Fraudulent
```

```
X = df.values
y = df['Class'].values
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.15, random_st
```

```
print("X =", X.size)
print("Y =", y.size)
print("X_train=", X_train.size)
print("X_test=", X_test.size)
print("y_train=", y_train.size)
print("y_test=", y_test.size)
```

```
X = 8829017
Y = 284807
X_train= 7504635
X_test= 1324382
```

```
y_train= 242085
y_test= 42722

res = []
tree_max_n = 2
tree_max_acc = 0
tree_max_f1 = 0
tree_knn_yhat = None
for i in range(1,10):
    tree_model = DecisionTreeClassifier(max_depth = i, criterion = 'entropy')
    tree_model.fit(X_train, y_train)
    tree_yhat_loc = tree_model.predict(X_test)
    tree_acc = accuracy_score(y_test, tree_yhat_loc)
    tree_f1 = f1_score(y_test, tree_yhat_loc)
    if tree_acc > tree_max_acc and tree_f1 > tree_max_f1:
        tree_max_acc = tree_acc
        tree_knn_yhat = tree_yhat_loc
        tree_max_f1 = tree_f1
        tree_max_n = i

print("max acc", tree_max_acc)
print("max f1", tree_max_f1)
print("max n", tree_max_n)
# max index 0

max acc 0.9994557775359011
max f1 0.835978835978836
max n 6

xgb = XGBClassifier(max_depth = 4)
xgb.fit(X_train, y_train)
xgb_yhat = xgb.predict(X_test)

print(accuracy_score(y_test, xgb_yhat))
print(f1_score(y_test, xgb_yhat))

0.9994733330992591
0.8421052631578948

rf = RandomForestClassifier(max_depth = 4)
rf.fit(X_train, y_train)
rf_yhat = rf.predict(X_test)
rf_max_n = 0
rf_max_acc = 0
rf_max_f1 = 0
rf_knn_yhat = None
for i in range(2,10):
    rf = RandomForestClassifier(max_depth = i)
    rf.fit(X_train, y_train)
    rf_yhat_l = rf.predict(X_test)
```

```

rf_acc = accuracy_score(y_test, knn_yhat_l)
rf_f1 = f1_score(y_test, knn_yhat_l)
if rf_acc > rf_max_acc and rf_f1 > rf_max_f1:
    rf_max_acc = rf_acc
    rf_yhat = rf_yhat_l
    rf_max_f1 = rf_f1
    rf_max_n = i

```

```

clg = GaussianNB()
clg.fit(X_train, y_train)
clg_yhat = clg.predict(X_test)
print(accuracy_score(y_test, clg_yhat))
print(f1_score(y_test, clg_yhat))

```

```

0.9783364348161933
0.12233285917496445

```

```

res = []
knn_max_n = 0
knn_max_acc = 0
knn_max_f1 = 0
knn_knn_yhat = None
for i in range(2,10):
    knn = KNeighborsClassifier(n_neighbors = i)
    knn.fit(X_train, y_train)
    knn_yhat_l = knn.predict(X_test)
    res.append(knn_yhat_l)
    knn_acc = accuracy_score(y_test, knn_yhat_l)
    knn_f1 = f1_score(y_test, knn_yhat_l)
    if knn_acc > knn_max_acc and knn_f1 > knn_max_f1:
        knn_max_acc = knn_acc
        knn_yhat = knn_yhat_l
        knn_max_f1 = knn_f1
        knn_max_n = i

print("max acc", knn_max_acc)
print("max f1", knn_max_f1)
print("max n", knn_max_n)
# max index 0

max acc 0.999403110845827
max f1 0.8089887640449437
max n 2

```

```

for el in res:
    print(accuracy_score(y_test, el))
    print(f1_score(y_test, el))

```

```

print(f1_score(y_test, e1))

0.999403110845827
0.8089887640449437
0.999385555282469
0.8066298342541437
0.9993504441557529
0.7909604519774012
0.9993328885923949
0.7865168539325842
0.9992977774656788
0.7701149425287357
0.9993153330290369
0.7771428571428572
0.9992451107756047
0.7485380116959064
0.9992451107756047
0.7485380116959064

lr = LogisticRegression()
lr.fit(X_train, y_train)
lr_yhat = lr.predict(X_test)
print(accuracy_score(y_test, lr_yhat))
print(f1_score(y_test, lr_yhat))

0.9992451107756047
0.7542857142857143
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py:940: Cc
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)

svm = SVC()
svm.fit(X_train, y_train)
svm_yhat = svm.predict(X_test)

print(accuracy_score(y_test, svm_yhat))
print(f1_score(y_test, svm_yhat))

0.998735999438222
0.5

def plot_confusion_matrix(cm, classes, title, normalize = False, cmap = plt.cm.Blues)
    title = 'Confusion Matrix of {}'.format(title)
    if normalize:
        cm = cm.astype(float) / cm.sum(axis=1)[:, np.newaxis]

    plt.imshow(cm, interpolation = 'nearest', cmap = cmap)
    plt.title(title)

```

```

plt.title(title)
plt.colorbar()
tick_marks = np.arange(len(classes))
plt.xticks(tick_marks, classes, rotation = 45)
plt.yticks(tick_marks, classes)

fmt = '.2f' if normalize else 'd'
thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, format(cm[i, j], fmt),
             horizontalalignment = 'center',
             color = 'white' if cm[i, j] > thresh else 'black')

plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')

```

```

file = tf.keras.utils
raw_df = pd.read_csv('creditcard 2.csv')
raw_df.head()

```

	Time	V1	V2	V3	V4	V5	V6	V7	V8
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098658
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085106
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247675
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377435
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533

```

cleaned_df = raw_df.copy()
cleaned_df.pop('Time')
eps = 0.001
cleaned_df['Log Ammount'] = np.log(cleaned_df.pop('Amount')+eps)

```

```

train_df, test_df = train_test_split(cleaned_df, test_size=0.2)
train_df, val_df = train_test_split(train_df, test_size=0.2)

```

```

train_labels = np.array(train_df.pop('Class'))
bool_train_labels = train_labels != 0
val_labels = np.array(val_df.pop('Class'))
test_labels = np.array(test_df.pop('Class'))

```

```

train_features = np.array(train_df)
val_features = np.array(val_df)
test_features = np.array(test_df)

```

```

scaler = StandardScaler()
train_features = scaler.fit_transform(train_features)

val_features = scaler.transform(val_features)
test_features = scaler.transform(test_features)

train_features = np.clip(train_features, -5, 5)
val_features = np.clip(val_features, -5, 5)
test_features = np.clip(test_features, -5, 5)

print('Training labels shape:', train_labels.shape)
print('Validation labels shape:', val_labels.shape)
print('Test labels shape:', test_labels.shape)

print('Training features shape:', train_features.shape)
print('Validation features shape:', val_features.shape)
print('Test features shape:', test_features.shape)

Training labels shape: (182276,)
Validation labels shape: (45569,)
Test labels shape: (56962,)
Training features shape: (182276, 29)
Validation features shape: (45569, 29)
Test features shape: (56962, 29)

METRICS = [
    keras.metrics.TruePositives(name='tp'),
    keras.metrics.FalsePositives(name='fp'),
    keras.metrics.TrueNegatives(name='tn'),
    keras.metrics.FalseNegatives(name='fn'),
    keras.metrics.BinaryAccuracy(name='accuracy'),
    keras.metrics.Precision(name='precision'),
    keras.metrics.Recall(name='recall'),
    keras.metrics.AUC(name='auc'),
    keras.metrics.AUC(name='prc', curve='PR'), # precision-recall curve
]

def make_model(metrics=METRICS, output_bias=None):
    if output_bias is not None:
        output_bias = tf.keras.initializers.Constant(output_bias)
    model = keras.Sequential([
        keras.layers.Dense(
            16, activation='relu',
            input_shape=(train_features.shape[-1],)),
        keras.layers.Dropout(0.5),
        keras.layers.Dense(1, activation='sigmoid',
                           bias_initializer=output_bias),
    ])

```

```

model.compile(
    optimizer=keras.optimizers.Adam(lr=1e-3),
    loss=keras.losses.BinaryCrossentropy(),
    metrics=metrics)

```

```

return model

```

```

EPOCHS = 100

```

```

BATCH_SIZE = 2048

```

```

early_stopping = tf.keras.callbacks.EarlyStopping(
    monitor='val_prc',
    verbose=1,
    patience=10,
    mode='max',
    restore_best_weights=True)

```

```

model = make_model()

```

```

model.summary()

```

```

model.predict(train_features[:10])

```

```

initial_weights = os.path.join(tempfile.mkdtemp(), 'initial_weights')

```

```

model.save_weights(initial_weights)

```

```

Model: "sequential_2"

```

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 16)	480
dropout_2 (Dropout)	(None, 16)	0
dense_5 (Dense)	(None, 1)	17
Total params: 497		
Trainable params: 497		
Non-trainable params: 0		

```

model = make_model()

```

```

model.load_weights(initial_weights)

```

```

baseline_history = model.fit(
    train_features,
    train_labels,
    batch_size=BATCH_SIZE,
    epochs=EPOCHS,
    callbacks=[early_stopping],
    validation_data=(val_features, val_labels))

```

```

Epoch 7/100

```

```

90/90 [=====] - 1s 11ms/step - loss: 0.0308 - tp: 92.68

```

```

Epoch 8/100

```

```

90/90 [=====] - 1s 10ms/step - loss: 0.0260 - tp: 92.31

```



```

Epoch 9/100
90/90 [=====] - 1s 10ms/step - loss: 0.0229 - tp: 102.8
Epoch 10/100
90/90 [=====] - 1s 10ms/step - loss: 0.0213 - tp: 97.14
Epoch 11/100
90/90 [=====] - 1s 10ms/step - loss: 0.0191 - tp: 106.2
Epoch 12/100
90/90 [=====] - 1s 10ms/step - loss: 0.0173 - tp: 106.9
Epoch 13/100
90/90 [=====] - 1s 10ms/step - loss: 0.0171 - tp: 103.1
Epoch 14/100
90/90 [=====] - 1s 11ms/step - loss: 0.0145 - tp: 112.5
Epoch 15/100
90/90 [=====] - 1s 11ms/step - loss: 0.0142 - tp: 107.4
Epoch 16/100
90/90 [=====] - 1s 11ms/step - loss: 0.0134 - tp: 115.6
Epoch 17/100
90/90 [=====] - 1s 11ms/step - loss: 0.0127 - tp: 102.3
Epoch 18/100
90/90 [=====] - 1s 11ms/step - loss: 0.0119 - tp: 105.0
Epoch 19/100
90/90 [=====] - 1s 11ms/step - loss: 0.0117 - tp: 112.1
Epoch 20/100
90/90 [=====] - 1s 11ms/step - loss: 0.0101 - tp: 105.7
Epoch 21/100
90/90 [=====] - 1s 11ms/step - loss: 0.0097 - tp: 119.4
Epoch 22/100
90/90 [=====] - 1s 11ms/step - loss: 0.0101 - tp: 116.6
Epoch 23/100
90/90 [=====] - 1s 10ms/step - loss: 0.0100 - tp: 106.1
Epoch 24/100
90/90 [=====] - 1s 11ms/step - loss: 0.0096 - tp: 103.1
Epoch 25/100
90/90 [=====] - 1s 11ms/step - loss: 0.0088 - tp: 123.1
Epoch 26/100
90/90 [=====] - 1s 11ms/step - loss: 0.0078 - tp: 101.5
Epoch 27/100
90/90 [=====] - 1s 10ms/step - loss: 0.0086 - tp: 109.0
Epoch 28/100
90/90 [=====] - 1s 11ms/step - loss: 0.0076 - tp: 121.2
Epoch 29/100
90/90 [=====] - 1s 11ms/step - loss: 0.0078 - tp: 115.5
Epoch 30/100
90/90 [=====] - 1s 11ms/step - loss: 0.0079 - tp: 112.0
Epoch 31/100
90/90 [=====] - 1s 10ms/step - loss: 0.0077 - tp: 105.7
Epoch 32/100
90/90 [=====] - 1s 11ms/step - loss: 0.0073 - tp: 107.7
Epoch 33/100
90/90 [=====] - 1s 11ms/step - loss: 0.0066 - tp: 112.2
Epoch 34/100
90/90 [=====] - 1s 11ms/step - loss: 0.0072 - tp: 110.4
Epoch 35/100
90/90 [=====] - 1s 11ms/step - loss: 0.0067 - tp: 113.9
Epoch 36/100

```





```
def plot_metrics(history):
```

```

metrics = ['loss', 'prc', 'precision', 'recall']
for n, metric in enumerate(metrics):
    name = metric.replace("_", " ").capitalize()
    plt.subplot(2,2,n+1)
    plt.plot(history.epoch, history.history[metric], color=colors[0], label='Train')
    plt.plot(history.epoch, history.history['val_'+metric],
              color=colors[0], linestyle="--", label='Val')
    plt.xlabel('Epoch')
    plt.ylabel(name)
    if metric == 'loss':
        plt.ylim([0, plt.ylim()[1]])
    elif metric == 'auc':
        plt.ylim([0.8,1])
    else:
        plt.ylim([0,1])

    plt.legend()
plot_metrics(baseline_history)

```



```
train_predictions_baseline = model.predict(train_features, batch_size=BATCH_SIZE)
test_predictions_baseline = model.predict(test_features, batch_size=BATCH_SIZE)
def plot_cm(labels, predictions, p=0.5):
    cm = confusion_matrix(labels, predictions > p)
    plt.figure(figsize=(5,5))
    sns.heatmap(cm, annot=True, fmt="d")
    plt.title('Confusion matrix @{:.2f}'.format(p))
    plt.ylabel('Actual label')
    plt.xlabel('Predicted label')

    print('Legitimate Transactions Detected (True Negatives): ', cm[0][0])
    print('Legitimate Transactions Incorrectly Detected (False Positives): ', cm[0][1])
    print('Fraudulent Transactions Missed (False Negatives): ', cm[1][0])
    print('Fraudulent Transactions Detected (True Positives): ', cm[1][1])
    print('Total Fraudulent Transactions: ', np.sum(cm[1]))

baseline_results = model.evaluate(test_features, test_labels,
                                  batch_size=BATCH_SIZE, verbose=0)
for name, value in zip(model.metrics_names, baseline_results):
    print(name, ': ', value)
print()

plot_cm(test_labels, test_predictions_baseline)
```

```
loss : 0.004249707330018282
tp : 75.0
fp : 13.0
tn : 56848.0
fn : 26.0
accuracy : 0.9993153214454651

neg, pos = np.bincount(raw_df['Class'])
total = neg + pos
weight_for_0 = (1 / neg)*(total)/2.0
weight_for_1 = (1 / pos)*(total)/2.0

class_weight = {0: weight_for_0, 1: weight_for_1}

print('Weight for class 0: {:.2f}'.format(weight_for_0))
print('Weight for class 1: {:.2f}'.format(weight_for_1))
weighted_model = make_model()
weighted_model.load_weights(initial_weights)






weighted_history = weighted_model.fit(
    train_features,
    train_labels,
    batch_size=BATCH_SIZE,
    epochs=EPOCHS,
    callbacks=[early_stopping],
    validation_data=(val_features, val_labels),
    class_weight=class_weight)


plot_metrics(weighted_history)
train_predictions_weighted = weighted_model.predict(train_features, batch_size=BATCH_SIZE)
test_predictions_weighted = weighted_model.predict(test_features, batch_size=BATCH_SIZE)
weighted_results = weighted_model.evaluate(test_features, test_labels,
                                           batch_size=BATCH_SIZE, verbose=0)
for name, value in zip(weighted_model.metrics_names, weighted_results):
    print(name, ': ', value)
print()

plot_cm(test_labels, test_predictions_weighted)
```

```
-----  
auc : 0.9494603276252747  
prc : 0.685690701007843
```

```
Legitimate Transactions Detected (True Negatives): 56133  
Legitimate Transactions Incorrectly Detected (False Positives): 728  
Fraudulent Transactions Missed (False Negatives): 16  
Fraudulent Transactions Detected (True Positives): 85  
Total Fraudulent Transactions: 101
```





```
print("Accuracy of XG Boost =", accuracy_score(y_test, xgb_yhat), "with F1 Score =",
print("Accuracy of CLG =", accuracy_score(y_test, clg_yhat), "with F1 Score =", f1_sc
print("Accuracy of NKK =", accuracy_score(y_test, knn_yhat), "with F1 Score =", f1_sc
print("Accuracy of Lin. Reg =", accuracy_score(y_test, lr_yhat), "with F1 Score =", f
print("Accuracy of SVM =", accuracy_score(y_test, svm_yhat), "with F1 Score =", f1_sc
print("Accuracy of Decision Tree =", accuracy_score(y_test, tree_knn_yhat), "with F1
print("Accuracy of Random Forest =", accuracy_score(y_test, rf_yhat), "with F1 Score
```

```
Accuracy of XG Boost = 0.9994733330992591 with F1 Score = 0.8421052631578948
Accuracy of CLG = 0.9783364348161933 with F1 Score = 0.12233285917496445
```

Accuracy of NKK = 0.999403110845827 with F1 Score = 0.8089887640449437
Accuracy of Lin. Reg = 0.9992451107756047 with F1 Score = 0.7542857142857143
Accuracy of SVM = 0.998735999438222 with F1 Score = 0.5
Accuracy of Decision Tree = 0.9994557775359011 with F1 Score = 0.835978835978836
Accuracy of Random Forest = 0.9990519995786665 with F1 Score = 0.674698795180723

```
xgb_matrix = confusion_matrix(y_test, xgb_yhat, labels = [0, 1])
clg_matrix = confusion_matrix(y_test, clg_yhat, labels = [0, 1])
knn_matrix = confusion_matrix(y_test, knn_yhat, labels = [0, 1])
lr_matrix = confusion_matrix(y_test, lr_yhat, labels = [0, 1])
svm_matrix = confusion_matrix(y_test, svm_yhat, labels = [0, 1])
tree_matrix = confusion_matrix(y_test, tree_knn_yhat, labels = [0, 1])
rf_matrix = confusion_matrix(y_test, rf_yhat, labels = [0, 1])

plt.rcParams['figure.figsize'] = (6, 6)

knn_cm_plot = plot_confusion_matrix(xgb_matrix,
                                    classes = ['Non-Default(0)', 'Default(1)'],
                                    normalize = False, title = 'XG Boost')

plt.show()
CLG_plot = plot_confusion_matrix(clg_matrix,
                                 classes = ['Non-Default(0)', 'Default(1)'],
                                 normalize = False, title = 'CLG')

plt.show()
knn_cm_plot = plot_confusion_matrix(knn_matrix,
                                    classes = ['Non-Default(0)', 'Default(1)'],
                                    normalize = False, title = '5/K-NN')

plt.show()
lreg_cm_plot = plot_confusion_matrix(lr_matrix,
                                     classes = ['Non-Default(0)', 'Default(1)'],
                                     normalize = False, title = 'Linear Reg.')

plt.show()
svm_cm_plot = plot_confusion_matrix(svm_matrix,
                                    classes = ['Non-Default(0)', 'Default(1)'],
                                    normalize = False, title = 'SVM')

plt.show()
dt_cm_plot = plot_confusion_matrix(tree_matrix,
                                    classes = ['Non-Default(0)', 'Default(1)'],
                                    normalize = False, title = 'Decision Tree')

plt.show()
rf_cm_plot = plot_confusion_matrix(rf_matrix,
                                   classes = ['Non-Default(0)', 'Default(1)'],
                                   normalize = False, title = 'Random Forest')

plt.show()
```

